

CSC3002 – Networks Assignment 1 – 2020

Socket programming project

Department of Computer Science
University of Cape Town, South Africa

February 18, 2020

This assignment is on networked applications where you will learn the basics of protocol design and socket programming for TCP connections in Java: how to create a socket, bind it to a specific address and port, as well as send and receive messages/files. You will develop a client-server application in groups of three students and this document describes the context/requirements (Section 1), what to submit (Section 3), and some basic information about socket programming (Section A).

1. Application Description

There are a number of file sharing protocols, many of them employing what is commonly referred to as P2P or Peer-to-Peer sharing. Usually, sharing files is done directly between two nodes in a network. A different approach for sharing is to use the client-server model where clients upload and download files from a shared server.

In this assignment, you are required to design and implement a client-server file sharing application that makes use of TCP sockets. You first need to implement a server that can receive and send files to clients. The client should be able to upload files (one at a time) to the server, and also to query the server for a list of files available. The client should also be able to request (download) a file from the server. As this is a network application, the clients and the server should be able to run on different hosts. The client interface can be in terminal, with appropriate user commands.

Your web server should accept and parse a request, get the requested file from the server's file system, create a response message consisting of the requested file preceded by header lines, and then send the response directly to the client. If the requested file is not present in the server, the server should send a "404 Not Found" message back to the client. You may start by developing a web server that handles only one request at a time. Once you have the basic application working, you should modify/upgrade to a multithreaded server that is capable of serving multiple requests simultaneously. Using threading, first create a main thread in which your server listens for clients at a fixed

port. When it receives a TCP connection request from a client, it will set up the TCP connection through another port and service the client request in a separate thread. There will be a separate TCP connection in a separate thread for each request/response pair.

You should also design your client to connect to the server using a TCP connection, send a request to the server, and display the server response as output. The client should take command line arguments specifying the server IP address, the port at which the server is listening, and the request, e.g requesting to download a particular file stored at the server.

2. Protocol Design and Specification

Messages in the application protocol could be either text, consisting of readable character strings (e.g ‘UPLOAD’), or binary format, where messages are blocks of structured binary data (eg. ‘1000001’ could be used to mean ‘UPLOAD’). Text-based protocols have the advantage of being human readable, hence provide for easier understanding, monitoring and testing. You are required to use text-based messaging in this application protocol design.

An important step in protocol design involves defining the framework of communication for the intended application. This entails specifying requirements and constraints such as: whether real-time interaction is expected; reliability (ie if we need to verify/check that every message is delivered correctly); and authentication/confidentiality issues, among others. In this assignment, you are to design the application that takes into consideration privacy/confidentiality, by allowing users indicate file sharing permissions, e.g should a file be visible/downloadable by anyone querying? Or, should a file be visible/downloadable (or not) to certain specific clients? Or can a shared secret-key be used to access/download files? The application protocol should be designed to include some of these features.

Specification of protocol messages involves defining the types and structure of messages. Three types of messages can be defined; commands, data transfer, and control. Command messages define the different stages of communication between parties, such as the initiation or termination of communication. Data transfer messages are used to carry the data that is exchanged between parties, and such data could be fragmented into several messages. Control messages manage the dialogue between parties, including such aspects as message acknowledgements, and retransmission requests.

The message structure constitutes at least the header and body. The header, whose structure must be known to the receiver, may contain fields that describe the actual data in the message. Some of the fields/information contained in the header might include the message type, the command, recipient information, and sequence information. The header generally has fixed size and contains clues that should help the receiver to understand the rest of the message.

The last aspect of the protocol design will be the communication rules that specify the sequence of messages at every stage of communication. This requires clearly specifying

messages and reactions for every communication scenario. You will need to represent such rules with sequence diagrams (at least two sequence diagrams will be required, one for upload process, and another for download process)

3. What you need to submit

You will be required to submit the following:

1. Your code with proper inline documentation (comments)
2. A report (about 4 pages) on the design and functionality of your chat application.
3. In addition, the report needs to include:
 - a) A list of features with a brief explanation for their inclusion
 - b) A protocol specification, detailing the message formats and structure. You are required to include sequence diagram(s).
 - c) Screenshots of the application revealing its features.
4. Oral presentation to be scheduled with the TAs and Tutors (oral to be done on day after submission deadline)

A. Multi-threaded Client/Server Applications—Sockets Programming in Java

A.1. What is a socket?

A socket is the one end-point of a two-way communication link between two programs running over the network. Running over the network means that the programs run on different computers, usually referred as the local and the remote computers. However one can run the two programs on the same computer. Such communicating programs constitutes a client/server application. The server implements a dedicated logic, called **service**. The clients connect to the server to get served, for example, to obtain some data or to ask for the computation of some data. Different client/server applications implement different kind of services.

To distinguish different services, a numbering convention was proposed. This convention uses integer numbers, called port numbers, to denote the services. A server implementing a service assigns a specific port number to the entry point of the service. There are no specific physical entry points for the services in a computer. The port numbers for services are stored in configuration files and are used by the computer software to create network connections.

A socket is a complex data structure that contains an internet address and a port number. A socket, however, is referenced by its descriptor, like a file which is referenced by a file descriptor. That is why, the sockets are accessed via an application programming interface (API) similar to the file input/output API. This makes the programming of network applications very simple. The two-way communication link between the two programs running on different computers is done by reading from and writing to the sockets created on these computers. The data read from a socket is the data wrote into the other socket of the link. And reciprocally, the the data wrote into a socket in the data read from the other socket of the link. These two sockets are created and linked during the connection creation phase. The link between two sockets is like a pipe that is implemented using a stack of protocols. This linking of the sockets involves that internally a socket has a much more complex data structure, or more precisely, a collaboration of data structures. Thus, a socket data structure is more than just an internet address and a port number. You have to imagine a socket as a data structure that contains at least the internet address and the port number on the local computer, and the internet address and the port number on the remote computer.

A.2. How is a network connection created?

A network connection is initiated by a client program when it creates a socket for the communication with the server. To create the socket in Java, the client calls the `Socket` constructor and passes the server address and the the specific server port number to it. At this stage the server must be started on the machine having the specified address and listening for connections on its specific port number.

The server uses a specific port dedicated only to listening for connection requests from

clients. It can not use this specific port for data communication with the clients because the server must be able to accept the client connection at any instant. So, its specific port is dedicated only to listening for new connection requests. The server side socket associated with specific port is called server socket. When a connection request arrives on this socket from the client side, the client and the server establish a connection. This connection is established as follows:

1. When the server receives a connection request on its specific server port, it creates a new socket for it and binds a port number to it.
2. It sends the new port number to the client to inform it that the connection is established.
3. The server goes on now by listening on two ports:
 - it waits for new incoming connection requests on its specific port, and
 - it reads and writes messages on established connection (on new port) with the accepted client.

The server communicates with the client by reading from and writing to the new port. If other connection requests arrive, the server accepts them in the similar way creating a new port for each new connection. Thus, at any instant, the server must be able to communicate simultaneously with many clients and to wait on the same time for incoming requests on its specific server port. The communication with each client is done via the sockets created for each communication.

The `java.net` package in the Java development environment provides the class `Socket` that implements the client side and the class `serverSocket` class that implements the server side sockets. The client and the server must agree on a protocol. They must agree on the language of the information transferred back and forth through the socket. There are two communication protocols:

- stream communication protocol
- datagram communication protocol

The stream communication protocol is known as TCP (transfer control protocol). TCP is a connection-oriented protocol. In order to communicate over the TCP protocol, a connection must first be established between two sockets. While one of the sockets listens for a connection request (server), the other asks for a connection (client). Once the two sockets are connected, they can be used to transmit and/or to receive data. When we say "two sockets are connected" we mean the fact that the server accepted a connection. As it was explained above the server creates a new local socket for the new connection. The process of the new local socket creation, however, is transparent for the client.

The datagram communication protocol, known as UDP (user datagram protocol), is a connectionless protocol. No connection is established before sending the data. The data

are sent in a packet called datagram. The datagram is sent like a request for establishing a connection. However, the datagram contains not only the addresses, it contains the user data also. Once it arrives to the destination the user data are read by the remote application and no connection is established. This protocol requires that each time a datagram is sent, the local socket and the remote socket addresses must also be sent in the datagram. These addresses are sent in each datagram.

The **java.net** package in the Java development environment provides the class DatagramSocket for programming datagram communications. UDP is an unreliable protocol. There is no guarantee that the datagrams will be delivered in a good order to the destination socket. For, example, a long text, split in several pages and sent one page per datagram, can be received in a different page order. On the other side, TCP is a reliable protocol. TCP guarantee that the pages will be received in the order in which they are sent. When programming TCP and UDP based applications in Java, different types of sockets are used. These sockets are implemented in different classes. The classes ServerSocket and Socket implement TCP based sockets and the class DatagramSocket implements UDP based sockets as follows:

- Stream socket to listen for client requests (TCP): the class ServerSocket.
- Stream socket (TCP): the class Socket.
- Datagram socket (UDP): the class DatagramSocket.

This document shows how to program TCP based client/server applications. The UDP oriented programming is not covered in document.

A.2.1. Opening a socket

The client side When programming a client, a socket must be opened like below:

```
Socket MyClient ;  
MyClient = new Socket("MachineName", PortNumber);
```

This code, however, must be put in a try/catch block to catch the IOException:

```
Socket MyClient ;  
try {  
    MyClient = new Socket("MachineName", PortNumber);  
}  
catch ( IOException e ) {  
    System . out . p r i n t l n ( e ) ;  
}
```

where:

- **MachineName** is the machine name to open a connection to and
- **PortNumber** is the port number on which the server to connect to is listening.

When selecting a port number, one has to keep in mind that the port numbers in the range from 0 to 1023 are reserved for standard services, such as email, FTP, HTTP, etc. For our service (the chat server) the port number should be chosen greater than 1023.

The server side: When programming a server, a server socket must be created first, like below:

```
ServerSocket MyService ;
try {
    MyService = new ServerSocket(PortNumber);
}
catch ( IOException e ) {
    System . out . p r i n t l n ( e ) ;
}
```

The server socket is dedicated to listen to and accept connections from clients. After accepting a request from a client the server creates a client socket to communicate (to send/receive data) with the client, like below :

```
Socket clientSocket = null ;
try {
    serviceSocket = MyService . accept () ;
}
catch ( IOException e ) {
    System . out . p r i n t l n ( e ) ;
}
```

Now the server can send/receive data to/from the clients. Since the sockets are like the file descriptors the send/receive operations are implemented like read/write file operations on the in- put/output streams.

A.2.2. Creating an input stream

On the client side, you can use the `DataInputStream` class to create an input stream to receive responses from the server:

```
DataInputStream input ;
try {
    input = new DataInputStream ( MyClient . getInputStream ( ) ) ;
}
```

```
        catch ( IOException e ) {  
            System . out . p r i n t l n ( e ) ;  
        }  
    }
```

The class `DataInputStream` allows you to read lines of text and Java primitive data types in a portable way. It has several read methods such as `read`, `readChar`, `readInt`, `readDouble`, and `readLine`. One has to use whichever function depending on the type of data to receive from the server.

On the server side, the `DataInputStream` is used to receive inputs from the client:

```
DataInputStream input ;  
try {  
    input = new DataInputStream( serviceSocket . getInputStream ( ));  
}  
catch ( IOException e ) {  
    System . out . p r i n t l n ( e ) ;  
}  
}
```

A.2.3. Create an output stream

On the client side, an output stream must be created to send the data to the server socket using the class **`PrintStream`** or **`DataOutputStream`** of `java.io` package:

```
PrintStream output ;  
try {  
    output = new PrintStream ( MyClient . getOutputStream ( ) ) ;  
}  
catch ( IOException e ) {  
    System . out . p r i n t l n ( e ) ;  
}  
}
```

The class **`PrintStream`** implements the methods for displaying Java primitive data types values, like **`write`** and **`println`** methods. Also, one may want to use the **`DataOutputStream`**:

```
DataOutputStream output ;  
try {  
    output = new DataOutputStream ( MyClient . getOutputStream ( ) ) ;  
}  
catch ( IOException e ) {  
    System . out . p r i n t l n ( e ) ;  
}
```



```
}
```

The class `DataOutputStream` allows you to write Java primitive data types; many of its methods write a single Java primitive type to the output stream.

On the server side, one can use the class `PrintStream` to send data to the client.

```
PrintStream output ;
try {
    output = new PrintStream(serviceSocket.getOutputStream());
}
catch ( IOException e ) {
    System . out . p r i n t l n ( e ) ;
}
```

A.2.4. Closing sockets

Closing a socket is like closing a file. You have to close a socket when you do not need it any more. The output and the input streams must be closed as well but before closing the socket.

On the client side you have to close the input and the output streams and the socket like below:

```
try {
    output . close ();
    input . close ();
    MyClient . close ();
}
catch ( IOException e ) {
    System . out . p r i n t l n ( e ) ;
}
```

On the server you have to close the input and output streams and the two sockets as follows:

```
try {
    output . close ();
    input . close ();
    serviceSocket . close ();
    MyService . close ();
}
catch ( IOException e ) {
```

```
        System . out . p r i n t l n ( e ) ;  
    }
```

Usually, on the server side you need to close only the client socket after the client gets served. The server socket is kept open as long as the server is running. A new client can connect to the server on the server socket to establish a new connection, that is, a new client socket¹.

¹Credits: these notes in the appendix are based on <http://www.ase.md/~aursu/ClientServer-Threads.html> and chapter 12, section 4 of ‘Introduction to Programming Using Java’, 6th ed. (v6.0.3, January 2014), by David J. Eck.